



# towboot

a Multiboot-compatible bootloader for  
EFI-based x86 systems

Niklas Sombert

2023-07-06

## 1 Introduction

Goals

Motivation

## 2 Background

Technologies used

Alternatives

## 3 Implementation and Usage

Boot process

Configuration

Memory Management

Workarounds

Changes made to other software

## 4 Testing

Components

Bootloader

## 5 Conclusion

Testing

Summary

Further Improvements

Code

- 1 Introduction
  - Goals
  - Motivation

## Goals

- UEFI application
- Multiboot-compatible bootloader

## Motivation

- firmware of current devices follows the UEFI standard
- Multiboot as the interface between kernel and bootloader
- hhuOS
  - Multiboot-compatible
  - licensing reasons
  - whole operating system

## ② Background

### Technologies used

UEFI

Rust

Multiboot

x86

ELF

TOML

### Alternatives

# Technologies used: UEFI

---

- Unified Extensible Firmware Interface<sup>1</sup>
- Tianocore<sup>2</sup>
- Open Virtual Machine Firmware<sup>3</sup>

---

<sup>1</sup> *Unified Extensible Firmware Interface (UEFI) Specification, Version 2.8 Errata B, May 2020.* UEFI Forum, 2020. URL: <https://uefi.org/sites/default/files/resources/UEFI%20Spec%202.8B%20May%202020.pdf> (visited on 10/13/2020).

<sup>2</sup> *tianocore/edk2.* URL: <https://github.com/tianocore/edk2> (visited on 01/01/2021).

<sup>3</sup> *OVMF.* [tianocore/tianocore.github.io](https://github.com/tianocore/tianocore.github.io). URL: <https://github.com/tianocore/tianocore.github.io/wiki/OVMF> (visited on 01/01/2021).

(without a Compatibility Support Module)

- relocatable PE<sup>4</sup> files
- loaded by the firmware (or another application)
- same CPU mode as the firmware
- `BootXXXX` variables
  - EFI System Partition
- removable media
  - `\EFI\BOOT\BOOTarch.EFI`

---

<sup>4</sup>Portable Executable

- `efi_main` function
- System Table
- Boot and Runtime Services
- opaque handles
- protocols with UUIDs



- Loaded Image Protocol
- Simple File System Protocol / File Protocol
- Simple Text Input Protocol / Simple Text Output Protocol
- Graphics Output Protocol<sup>5</sup>
- many more possible

---

<sup>5</sup>*Replacing VGA, GOP implementation for UEFI.* UEFI Summer Plugfest. 2011. URL: [https://uefi.org/sites/default/files/resources/UPFS11\\_P4\\_UEFI\\_GOP\\_AMD.pdf](https://uefi.org/sites/default/files/resources/UPFS11_P4_UEFI_GOP_AMD.pdf) (visited on 10/25/2020).

```
1  #![no_std]
2  #![no_main]
3  use uefi::prelude::*;
4  use uefi::proto::loaded_image::LoadedImage;
5  #[entry]
6  fn efi_main(image: Handle, mut systab: SystemTable<Boot>) -> Status {
7      uefi_services::init(&mut systab).expect("Failed to initialize utilit
8      let loaded_image = systab.boot_services()
9          .open_protocol_exclusive::<LoadedImage>(image)
10         .expect("Failed to open loaded image protocol");
11      let load_options = loaded_image.load_options_as_cstr16()
12         .expect("Failed to get load options");
13      writeln!(systab.stdout(), "got load options: {}", load_options);
14      Status::SUCCESS
15 }
```

## Technologies used: Rust<sup>6</sup>

---

- system programming language
- low-level code, high-level abstractions
- type- and memory-safety (excluding `unsafe`)
- cargo

---

<sup>6</sup> *Rust Programming Language*. URL: <https://www.rust-lang.org/> (visited on 01/01/2021).

- `i686-unknown-uefi` and `x86_64-unknown-uefi` targets<sup>7</sup>
  - Tier 2
  - `no_std`
- `uefi` crate<sup>8</sup>

---

<sup>7</sup> *The rustc book*. Chap. Platform Support. URL: <https://doc.rust-lang.org/nightly/rustc/platform-support.html> (visited on 06/23/2023).

<sup>8</sup> Gabriel Majeri. *uefi*. crates.io. URL: <https://crates.io/crates/uefi> (visited on 12/04/2020).

# Technologies used: Multiboot<sup>1213</sup>

- standard for kernels and bootloaders
- from the GNU GRUB<sup>9</sup> project
  - GNU kernel: gnumach<sup>10</sup> (for GNU HURD<sup>11</sup>)
- two major versions

---

<sup>9</sup>GNU GRUB. GNU. URL: <https://www.gnu.org/software/grub/> (visited on 10/13/2020).

<sup>10</sup>gnumach. GNU. URL: <https://www.gnu.org/software/hurd/microkernel/mach/gnumach.html> (visited on 12/05/2020).

<sup>11</sup>GNU Hurd. GNU. URL: <https://www.gnu.org/software/hurd/> (visited on 01/01/2021).

<sup>12</sup>Bryan Ford and Erich Stefan Boleyn. *Multiboot Specification version 0.6.96*. Free Software Foundation, Inc. URL: <https://www.gnu.org/software/grub/manual/multiboot/multiboot.html> (visited on 10/13/2020).

<sup>13</sup>Bryan Ford and Erich Stefan Boleyn. *Multiboot2 Specification version 2.0*. GNU. Free Software Foundation, Inc. URL: <https://www.gnu.org/software/grub/manual/multiboot2/multiboot.html> (visited on 10/13/2020).

- load kernel and modules (ELF or flat binary)
  - ELF or flat binary
  - header: static information
    - module alignment
    - require passing memory or video information
    - preferred video mode
    - binary loading and entry
    - additional flags

- pass information about the system
  - struct, `EAX` and `EBX` registers
    - memory (upper, lower, map)
    - kernel command line
    - modules (and command line)
    - symbol
    - bootloader name
    - framebuffer
    - legacy stuff (boot device, drives, config table, APM, VBE)
    - on Multiboot 2: SMBIOS, ACPI, UEFI, network
  - machine state
    - 32-bit Protected Mode
    - no interrupts
    - no paging
    - no PAE (in practice)
    - Multiboot 2: also native UEFI mode

## Technologies used: Multiboot

---

- multiboot crate<sup>14</sup>, PR<sup>15</sup>
- multiboot2 crate<sup>16</sup>, PR<sup>17</sup>

---

<sup>14</sup>Gerd Zellweger. *multiboot*. crates.io. URL: <https://crates.io/crates/multiboot> (visited on 01/04/2021).

<sup>15</sup>Niklas Sombert. *Add some functionality for using this in bootloaders (Pull Request 8)*. gz/rust-multiboot. URL: <https://github.com/gz/rust-multiboot/pull/8> (visited on 01/04/2021).

<sup>16</sup>Philipp Oppermann. *multiboot2*. crates.io. URL: <https://crates.io/crates/multiboot2> (visited on 06/23/2023).

<sup>17</sup>Niklas Sombert. *Add a builder to multiboot2 (Pull Request 133)*. rust-osdev/multiboot2. URL: <https://github.com/rust-osdev/multiboot2/pull/133> (visited on 06/23/2023).



- i686 (32-bit) and x86\_64 (64-bit)
- multiple modes<sup>18</sup>
  - Long Mode
    - 64-Bit Mode
    - Compatibility Mode (32-Bit)
  - Legacy Mode
    - Protected Mode (32-Bit or 16-Bit)
    - Virtual-8086 Mode (16-Bit)
    - Real Mode (16-Bit)
- Alternatives:
  - MIPS (Multiboot 2)
  - Itanium, ARM, RISC-V (UEFI)

---

<sup>18</sup>AMD64 Architecture Programmer's Manual. 2020. Chap. 1.2 Modes of Operation. URL: <https://www.amd.com/system/files/TechDocs/40332.pdf> (visited on 01/04/2021).

- Multiboot machine state requires many x86-specifics
  - native mode to 32-bit Protected Mode
  - may have to switch through Compatibility Mode
- memory
  - Real Mode (8086): 16-bit pointers, 24-bit address bus, 1MB, memory hole, lower 640KB
  - 286: 24-bit address bus: 1MB to approx 10MB
  - i686: 4GB virtually, physically (more with PAE)
  - x86\_64: physically 256TB memory
  - *segmentation*

- in Real Mode, segment registers are offsets

```
1      mov [ax], 0xab ; writes to *(ds * 16 + ax)
2      mov es:[ax], 0xab ; similar, but with es
3      jmp 0xabc ; actually jumps to cs * 16 + 0xabc ("near-jump")
```

- they're settable, usually

```
1      mov ds, 0x12
2      jmp 0xde:0xabc ; cs can only be set with a "far-jump"
```

- in Protected Mode, segment registers are indices into the Global Descriptor Table
- for a very simple OS with just one application it might look like this:

index	content
0	NULL
1	system code
2	system data
3	application code
4	application data

- information contained in these Segment Descriptors:
  - base, limit, access, flags
  - not layed out sequentially!
- `x86 crate`<sup>19</sup> has a builder

<sup>19</sup>Gerd Zellweger. `x86`. `crates.io`. URL: <https://crates.io/crates/x86> (visited on 06/30/2023).

```
1  let code_segment_builder = SegmentDescriptorBuilder::code_descriptor(  
2      0, u32::MAX, CodeSegmentType::ExecuteRead,  
3  );  
4  let code_segment: Descriptor = code_segment_builder  
5      .present().limit_granularity_4kb().db().finish();  
6  let data_segment_builder = SegmentDescriptorBuilder::data_descriptor(  
7      0, u32::MAX, DataSegmentType::ReadWrite,  
8  );  
9  let data_segment: Descriptor = data_segment_builder  
10     .present().limit_granularity_4kb().db().finish();  
11  let gdt = DescriptorTablePointer::new_from_slice(  
12     &[Descriptor::NULL, code_segment, data_segment]  
13  );
```

## Technologies used: x86: code (2/3)

```
1  unsafe { x86::irq::disable();
2          x86::dtables::lgdt(&gdt);
3          x86::dtables::lidt(&DescriptorTablePointer::default()); // invalid
4          asm!("push 0x08", // code segment
5              "lea rbx, [2f]",
6              "push rbx",
7              "retfq", // overwrite CS
8              "2:", // compatibility mode, yay.
9              ".code32",
10             "mov eax, 0x10", // data segment
11             "mov ds, eax",
12             "mov es, eax",
13             "mov fs, eax",
14             "mov gs, eax",
15             "mov ss, eax",
```

## Technologies used: x86: code (3/3)

```
1      "mov ecx, cr0",
2      "and ecx, ~(1<<31)", // disable paging
3      "or ecx, 1", // enable protected mode
4      "mov cr0, ecx",
5
6      "mov ecx, cr4",
7      "and ecx, ~(1<<5)", // disable PAE
8      "mov cr4, ecx",
9
10     "mov ecx, 0xC0000080",
11     "rdmsr",
12     "and eax, ~(1<<8)", // disable long mode
13     "wrmsr",
14 );
15 }
```

# Technologies used: ELF

---

- Executable and Linkable Format<sup>20</sup>
- support needed to be Multiboot-compliant
- header
- Program Header
- Section Header
- `goblin crate`<sup>21</sup>

---

<sup>20</sup>[Wikipedia contributors](https://en.wikipedia.org/w/index.php?title=Executable_and_Linkable_Format&oldid=992446183). *Executable and Linkable Format* — *Wikipedia, The Free Encyclopedia*. 2020. URL: [https://en.wikipedia.org/w/index.php?title=Executable\\_and\\_Linkable\\_Format&oldid=992446183](https://en.wikipedia.org/w/index.php?title=Executable_and_Linkable_Format&oldid=992446183) (visited on 01/04/2021).

<sup>21</sup>[m4b et al. \*goblin\*. crates.io](https://crates.io/crates/goblin). URL: <https://crates.io/crates/goblin> (visited on 01/04/2021).



## Technologies used: TOML

---

- Tom's Obvious Minimal Language<sup>22</sup>
- format for configuration files
- simple, yet type-safe and well-defined
- `toml` crate<sup>23</sup>, PR<sup>24</sup>

---

<sup>22</sup>Tom Preston-Werner. *TOML: Tom's Obvious Minimal Language*. URL: <https://toml.io/> (visited on 10/14/2020).

<sup>23</sup>Eric Huss and Alex Crichton. *toml*. [crates.io](https://crates.io/crates/toml). URL: <https://crates.io/crates/toml> (visited on 01/04/2021).

<sup>24</sup>Thom Chiovoloni. *Added nostd support (Pull Request 429)*. [alexcrichton/toml-rs](https://github.com/alexcrichton/toml-rs). URL: <https://github.com/alexcrichton/toml-rs/pull/429> (visited on 04/14/2022).

- GNU GRUB<sup>25</sup>
  - widely used bootloader
  - supports various platforms and operating systems
  - menu, can edit entries interactively
  - complex scripting language
  - de-facto reference implementation of Multiboot
- Syslinux<sup>26</sup>
  - very lightweight bootloader
  - supports multiple platforms and some operating systems
  - current release 2014
- Limine<sup>27</sup>
  - supports various platforms and operating systems
  - also has an own boot protocol

---

<sup>25</sup> GNU GRUB.

<sup>26</sup> The Syslinux Project. URL: <https://www.syslinux.org/> (visited on 01/02/2021).

<sup>27</sup> Limine. URL: <https://limine-bootloader.org/> (visited on 06/08/2023).

## 3 Implementation and Usage

- Boot process

- Configuration

  - NVRAM

  - Command line parameters

  - Configuration file

- Memory Management

  - Stack

  - Heap

  - Whole pages

- Workarounds

- Changes made to other software

  - multiboot crates

  - hhuOS

  - miniarg

  - multiboot12

- ① initialize the bootloader
- ② determine what to boot
  - ① parse the command line parameters
  - ② parse the configuration file
  - ③ display a menu
- ③ boot the operating system
  - ① load the kernel
  - ② load the modules
  - ③ configure the graphics output
  - ④ prepare the information to pass to the kernel
  - ⑤ deinitialize most parts of the bootloader
  - ⑥ jump to the kernel's entry point

- binary encoding: CBOR<sup>28</sup> or JSON<sup>29</sup> (for instance)
- EFI variable
- requires tooling
- bound to a system
- firmware might “forget” configuration

<sup>28</sup>Carsten Bormann. *Concise Binary Object Representation*. URL: <https://cbor.io> (visited on 10/14/2020).

<sup>29</sup>Douglas Crockford. *Introducing JSON*. URL: <https://www.json.org/> (visited on 10/14/2020).

```
FS0:\> towboot.efi -help
```

Usage:

-config Load the specified configuration file instead of the default one.  
-kernel Don't load a configuration file, instead boot the specified kernel.  
-logLevel Set the log level. (This only applies if '-kernel' is specified.)  
-module Load a module with the given args. Can be specified multiple times.  
-quirk Enable a specific quirk. (Only applies when loading a kernel.)  
-help Displays all available options and how to use them.  
-version Displays the version of towboot

- `towboot.efi -kernel "kernel.elf quiet" -module "module1.bin -verbose" -module "module2.bin -quiet"`
- `towboot.efi -config config.toml`
- file names or arguments may contain spaces or quotes
- `BootXXXX` variables
- shell or from another bootloader
- boot manager in the firmware usually not very user-friendly
- firmware might “forget” configuration

# Configuration: Configuration file

```
1 default = "entry1" 12
2 timeout = 10 13
3 # log_level defaults to 'info' 14
4 15
5 [entries] 16
6     [entries.entry1] 17
7         name = "Operating System" 18
8         image = "kernel.elf" 19
9         argv = "quiet" 20
10        quirks = ["KeepResolution"] 21
11
```

```
[[entries.entry1.modules]]
    image = "module1.bin"
    argv = "--verbose"
```

```
[[entries.entry1.modules]]
    image = "module2.bin"
    argv = "--quiet"
```

- easy to read and modify by hand
- most bootloaders do this
- `towboot.toml`, placed in top-level directory of the ESP

# Memory Management: Stack

- local variables
- works always
- size needs to be known at compile time
- used for most runtime data, for example `MultibootInfo`
- tracked by `rustc` at compile time



- `alloc::alloc::alloc` is bound to `allocate_pool`
- tracked both by `rustc` at compile time and by the firmware at runtime
- used for everything with a dynamic size and no further requirements

## Memory Management: Whole pages

- can be allocated to specific address (used for the kernel)
- custom `mem::Allocation wrapper (impl Drop)`
- tracked by the firmware at runtime
- also used for modules

- stack probes<sup>30</sup> broken with Rust 1.49 (fixed with 1.51, at least)
- initially no support for floating point operations (now just for x86\_64)
- `hacks`  module

---

<sup>30</sup>*probestack.rs*. *rust-lang/compiler-builtins*. URL: <https://github.com/rust-lang/compiler-builtins/blob/63ccaf11f08fb5d0b39cc33884c5a1a63f547ace/src/probestack.rs> (visited on 01/04/2021).

- `multiboot crate`<sup>31</sup> was designed for use in kernels
- added support for parsing the header, for setting values in `MultibootInfo` and tests<sup>32</sup>
- most difficult part: integration with Rust's and UEFI's memory management
- the same goes for `multiboot2`<sup>3334</sup>

---

<sup>31</sup>Zellweger, *multiboot*.

<sup>32</sup>Sombert, *Add some functionality for using this in bootloaders (Pull Request 8)*.

<sup>33</sup>Oppermann, *multiboot2*.

<sup>34</sup>Sombert, *Add a builder to multiboot2 (Pull Request 133)*.

# Changes made to other software: hhuOS

- kernel assumed the Multiboot structs to be before the kernel in memory
- PR<sup>35</sup> to copy them to BSS very early
- kernel also looks for the free memory after data passed by the bootloader

---

<sup>35</sup>Niklas Sombert. *Copy the Multiboot info struct recursively to BSS before enabling paging (Pull Request 19).* hhuOS/hhuOS. URL: <https://github.com/hhuOS/hhuOS/pull/19> (visited on 01/09/2021).

## Changes made to other software: miniarg

- parses strings such as `program -foo "bar baz"`
- `no_std-compatible`
- released to crates.io<sup>36</sup>
- not specific to UEFI or Multiboot

---

<sup>36</sup>Niklas Sombert. *miniarg*. crates.io. URL: <https://crates.io/crates/miniarg> (visited on 12/05/2020).

## Changes made to other software: multiboot12

---

- abstraction over `multiboot` and `multiboot2` crates
- currently just laying around on GitHub<sup>37</sup>

---

<sup>37</sup>Niklas Sombert. *multiboot12*. URL: <https://github.com/Ytvw1D/multiboot12> (visited on 06/30/2023).

## 4 Testing

- Components
- Bootloader



## Components

- unit tests for `multiboot`, `multiboot2` and `miniarg`
- not (yet?) for `multiboot12`

## Bootloader

- manual end-to-end test, could be automated?
- Multiboot specs<sup>3839</sup> contain example kernels
- had to be modified to print to serial
- booting a whole operating system
- virtual machine

<sup>38</sup>Ford and Boleyn, *Multiboot Specification version 0.6.96*.

<sup>39</sup>Ford and Boleyn, *Multiboot2 Specification version 2.0*.

## 5 Conclusion

Testing

Summary

Further Improvements

Code

- successfully tested with the test kernel, hhuOS<sup>40</sup>, GNU HURD<sup>41</sup>, Lemon OS<sup>42</sup>, HelenOS<sup>43</sup> and NetBSD<sup>44</sup>
- tests with FlingOS<sup>45</sup> and OpenIndiana<sup>46</sup> failed

---

<sup>40</sup> *hhuOS*. URL: <https://hhuos.github.io/> (visited on 01/01/2021).

<sup>41</sup> *GNU Hurd*.

<sup>42</sup> *Lemon OS*. URL: <https://lemonos.org/> (visited on 06/30/2023).

<sup>43</sup> *HelenOS*. URL: <http://www.helenos.org/> (visited on 06/30/2023).

<sup>44</sup> *NetBSD*. URL: <https://netbsd.org/> (visited on 06/30/2023).

<sup>45</sup> *Edward Nutting*. *FlingOS*. URL: <http://www.flingos.co.uk/> (visited on 01/01/2021).

<sup>46</sup> *OpenIndiana*. URL: <https://www.openindiana.org/> (visited on 06/30/2023).

## Conclusion: Summary

---

- UEFI's abstractions make writing a bootloader very straight-forward
- `uefi-rs` maps them to Rust
- Rust support is good, but could be better
- Multiboot requires x86-specifics

## Conclusion: Further Improvements

---

- Secure Boot<sup>47</sup>
  - enabled by default on most modern systems
  - changes to kernels and the kernel file format needed to be secure
- 64-bit, UEFI-unaware kernels
- other CPU architectures?
- load improvements
- file paths

---

<sup>47</sup>Richard Wilkins and Brian Richardson. *UEFI Secure Boot In Modern Computer Security Solutions*. 2013. URL: [https://uefi.org/sites/default/files/resources/UEFI\\_Secure\\_Boot\\_in\\_Modern\\_Computer\\_Security\\_Solutions\\_2019.pdf](https://uefi.org/sites/default/files/resources/UEFI_Secure_Boot_in_Modern_Computer_Security_Solutions_2019.pdf) (visited on 10/25/2020).

## hhuOS/towboot

a bootloader for Multiboot kernels on UEFI systems  
written in Rust



2

Contributors



1

Issue



5

Stars



0

Forks



<https://github.com/hhuOS/towboot>