



Rust Crash Course

- Basic syntax
- The borrow checker
- Practice problems
- Convenience types
- Slides from Rohan Kumar, Rahul Kumar, and Edward Zeng
 - Used in the course CS 162: Operating Systems and Systems Programming, Prof. John Kubiawicz at Berkely University, USA

Why Rust


- Memory safe!
- Fast executables
- Fewer runtime bugs

- The compiler will try to guess the type

```
let x: i32 = 2;  
let x = 2;
```

- Variables are **immutable** by default

```
// This won't compile  
let x = 2;  
x = 2;
```



```
// This is OK  
let mut x = 2;  
x = 2;
```

- Rust **references** are like pointers in C, but are:
 - Always valid (no pointers to freed memory)
 - Never null
- The Rust compiler checks these properties at **compile time**
- Reference with `&`, dereference with `*`

```
let x = 2;  
let ptr = &x;  
assert_eq!(*ptr, 2);
```

- By default, references are **immutable**.
- This won't work:

```
let mut x = 2;  
let ptr = &x;  
*ptr = 3; // Cannot mutate through an immutable reference
```

- If you want to modify a variable through a pointer, you must have a **mutable reference**

```
let mut x = 2;  
let ptr = &mut x;  
*ptr = 3; // This is OK
```

- To obtain a mutable reference, the variable itself must be mutable:

```
let x = 2;  
let ptr = &mut x; // Cannot mutably reference an immutable variable
```

- Obtaining a reference in Rust is called **borrowing**.

- Signed integers: `i8`, `i16`, `i32`, `i64`, `i128`, `isize`
- Unsigned integers: `u8`, `u16`, `u32`, `u64`, `u128`, `usize`
- Floating point: `f32`, `f64`
- Boolean: `boolean`
- Character: `char` (use single quotes, e.g. `'a'`)

■ Tuples:

```
// The type annotation is unnecessary  
let my_tuple: (i32, char, bool) = (162, 'X', true);
```

■ Arrays:

```
// Arrays have a fixed size, as indicated in the (optional) type annotation  
let nums: [i32, 5] = [1, 2, 3, 4, 5];  
let second = nums[1];
```

■ Structs:

```
struct Coordinate {  
    x: i32,  
    y: i32,  
}  
  
let point = Coordinate { x: 5, y: 3, };
```

- A function that squares the input:

```
fn square(x: i32) -> i32 {  
    return x * x;  
}
```

- Equivalent to:

```
fn square(x: i32) -> i32 {  
    x * x  
}
```

- If a function does not return a value, just omit the return type:

```
fn add_one(x: &mut i32) {  
    *x += 1;  
}
```

- A function that squares the input:

```
fn main() {  
    println!("Hello world!");  
}
```

- `println!` is a **macro**. The latter always have the exclamation mark as suffix.
- The compiler expands macros during compilation and macros will be replaced by „regular“ Rust source code.

- Rust has conditional if statements. No parentheses for the boolean expressions!

```
let x = 4;  
if x > 5 {  
    println!("Greater than 5");  
} else if x > 3 {  
    println!("Greater than 3");  
} else {  
    println!("x was not big enough");  
}
```

- All expressions can evaluate to a value.

```
let x = 4;
let message = if x > 5 {
    "Greater than 5"
} else if x > 3 {
    "Greater than 3"
} else {
    "x was not big enough"
};
println!("{}", message)
```

```
loop {  
    println!("stuck in a loop");  
}
```

■ Use break to exit

```
loop {  
    break;  
}
```

■ Loop expressions can evaluate to a value, just like any other expression:

```
let mut count = 0;  
let three = loop {  
    count += 1;  
    if count >= 3 {  
        break count;  
    }  
};
```

- Rust while loops are fairly straightforward:

```
let mut count = 5;  
while count > 0 {  
    count -= 1;  
}
```

- For loop can iterate over a collection

```
let muts = [0, 1, 2, 3, 4];  
for num in nums {  
    println!("{}", num);  
}  
// Prints 0 1 2 3 4
```

- Range notion:

```
for num in 0..5 {  
    println!("{}", num);  
}  
// Prints 0 1 2 3 4
```


- Useful when a type should have only a few possible values

```
enum Coin {  
    Head,  
    Tail,  
}  
  
let a = Coin::Head;  
let b = Coin::Tail;
```

- Each value in an enum is called a **variant**.

■ Enums can also store data!

```
enum OperatingSystem {  
    Mac,  
    Windows,  
    Linux,  
    Other(String),  
}  
  
let a = OperatingSystem::Linux;  
let b = OperatingSystem::Other("Redox OS");
```

- Rust **match expressions** are like C switch statements. However, they must *always be exhaustive*.
- What is wrong here?

```
let num = 162;  
  
match num {  
    160 => println!(160);  
    161 => println!(161);  
    168 => println!(168);  
}
```

- The match is not exhaustive! What if num was 164 or 10?

- This **is** valid:

```
let num = 162;  
  
match num {  
    160 => println!(160);  
    161 => println!(161);  
    168 => println!(168);  
    _  => println!("another case");  
}
```

- The underscore matches anything that was not already matched.
- Each pattern in the match statement is called a **match arm**.

- Matching is very useful in combination with enums. Match expressions can also evaluate to a value (just like any other expression).

```
enum OperatingSystem {  
    Mac,  
    Windows,  
    Linux,  
    Other(String),  
}  
  
fn os_name(os: OperatingSystem) -> String {  
    match os {  
        Mac      => "mac".to_string(),  
        Windows  => "windows".to_string(),  
        Linux    => "linuxc".to_string(),  
        Other(s) => s,  
    }  
}
```

- Suppose we have the following struct definition:

```
struct Vector {  
    x: f64,  
    y: f64,  
}
```

- We might want to add 2 Vectors elementwise. Here is one way to do that:

```
fn add(v1: Vector, v2: Vector) -> Vector {  
    x: v1.x + v2.x,  
    y: v1.y + v2.y,  
}
```

- We can also do the same thing using `impl` blocks, which define methods related to a struct (or enum).

```
impl Vector {  
    fn add(&self, other: Vector) -> Self {  
        x: self.x + other.x,  
        y: self.y + other.y,  
    }  
}  
  
let sum = v1.add(&v2);  
  
// Can also be called like this  
let sum = Vector::add(&v1, &v2);
```

- `Self` is a shorthand for the type of the `impl` block (in this case, `Vector`).

Arguments to `impl` block functions

- First argument is `&self`: the compiler will immutably borrow the object

```
let v1 = ...;  
v1.add(&v2)
```

- Is approximately equivalent to

```
let v1 = ...;  
let reference = &v1;  
Vector::add(reference, &v2);
```


Arguments to `impl` block functions

- Now first argument is `&mut self`: the compiler will mutably borrow the object

```
impl Vector {  
    fn double(&mut self) {  
        self.x *= 2;  
        self.y *= 2;  
    }  
}
```

```
let mut v1 = ...;  
v1.double();
```

- Is approximately equivalent to

```
let mut v1 = ...;  
let reference = &mut v1;  
Vector::double(reference);
```

The Borrow Checker

- A very important concept
- At the beginning sometimes difficult

Why have a borrow checker?

- We are trying to solve the problem of when to allocate and deallocate memory
- In C, you have to do this manually using `malloc` and `free`.
- In Java, a garbage collector runs periodically to free objects that are no longer usable
- In Rust, the borrow checker automatically determines when a value is unusable, and inserts code to free it at that point
- Automatic memory management without the overhead of garbage collection!
(Sometimes also called compile-time garbage collection)

- The **borrow checker** is what makes Rust *very* different from other languages
- The borrow checker verifies a set of rules at compile time. It does the magic of making sure your references are always valid.
- The borrow checker rules can initially seem mysterious. But they become easier with practice.
- We'll incrementally build up some of the basic borrow checking rules.

- Here's one set of borrow checking rules:
- Every value has one and only one owner
- When a value's owner goes out of **scope**, the value is **dropped** (= freed)
- To manually drop a value *v* *early*, call `drop(v)`
- Is it possible to drop a value late?
- No. A value is dropped when it goes out of scope. You can modify your code so that the scope is larger, but you cannot call `drop(v)` if *v* is not in scope.

- **Scopes** are enclosed in curly braces:

```
fn do_stuff() {  
    let a = String::from("hello");  
    {  
        let b = String::from("goodbye");  
        // Can access a and b  
        // ...  
        // b goes out of scope and is dropped  
    }  
    // Cannot access b here: it is out of scope  
    // a is dropped at the end of the function  
}
```

- Functions, loops, if statements etc. have their own scope. You can also create nested scopes using curly braces.

- Every value has one owner. Sometimes that owner can change. This is called a **move**
- Assignment moves values.
- This is invalid:

```
let s1 = String::from("my string");  
let s2 = s1; // Ownership of the string moves from s1 to s2.  
  
// s1 no longer owns the string, so we can't access data via s1.  
println!("{}", s1); // This is an error; data has been moved out of s1.
```

- This is fine:

```
let s1 = String::from("my string");  
let s2 = s1; // Ownership of the string moves from s1 to s2.  
println!("{}", s2); // This is okay; s2 owns the string now.
```

- If you need multiple variables to own data, you can clone a value:

```
let s1 = String::from("my string");  
let s2 = s1.clone(); // s2 is a clone of s1  
println!("{}", s2); // This is okay; s2 owns its data.  
println!("{}", s1); // This is okay; s1 also owns its data.
```

- Note that cloning is usually *expensive*. In this case, we are allocating memory for a string *twice*.
- In the previous examples, we only allocated space for the string once.
- Cloned values are completely independent of the value they were cloned from. If `s1` is modified, code using `s2` will not see those changes.

- There is one case in which a move behaves like a clone: when the type is Copy.
- Copy is a **trait** (ie. interface) that indicates that a value is copied whenever it is used.
- For example, integers are Copy:

```
let x = 5;  
let y = x;  
println!("{}", x, y);
```

- This is fine because the value in x (5) is copied, not moved. So x and y both own their values and can be accessed.

- In general, types that require heap allocations are not Copy.
- Copy: integers, floats, booleans, chars, immutable references, and compound types containing only Copy types.
- Not Copy: Strings, Vectors, mutable references, and compound types containing at least one non-Copy type.

- Consider the following struct:

```
struct Person {  
    id: u64,  
    age: u32,  
}
```

- Is it Copy?
- **No**. But shouldn't it be Copy, since it only contains Copy types?
- We must explicitly tell the Rust compiler we wish to make it Copy

```
#[derive(Copy)]  
struct Person {  
    id: u64,  
    age: u32,  
}
```

- Can also make structs cloneable by adding `#[derive(Clone)]`.

- Passing a value to a function moves the value:

```
fn main() {  
    let s = String::from("hello");  
    do_stuff(s);  
    // s no longer accessible; it was moved into do_stuff  
}  
  
fn do_stuff(s: String) {  
    // do stuff with s  
}
```

- Returning a value from a function moves the value to the caller:

```
fn main() {  
    let s = get_string();  
    // We can now use s, which owns the string  
}  
  
fn get_string() -> String {  
    String::from("hello")  
}
```

- (Aside: Rust functions generally don't take in Strings, but don't worry about this for now.)

- *You can have aliasing or mutability, but not both.*
- Aliasing:

```
let x = 162;  
let p1 = &x;  
let p2 = p1;
```

- x is **aliased**: multiple variables can read (not modify) the variable.

```
let mut x = 162;  
let p1 = &mut x;
```

- x is **mutable**; p1 can modify the contents of x.
- This is *forbidden*:

```
let mut x = 162;  
let p1 = &mut x;  
let p2 = &mut x;
```

No dangling pointers

- Rust ensures that you don't create dangling references at *compile time*.
- This code won't compile:

```
fn get_string() -> &String {  
    let s = String::from("hi");  
    &s  
}  
// s is dropped at the end of this function,  
// so &s would be a dangling pointer.  
// The Rust compiler won't allow this.
```

- Key point: a reference can never outlive the value it points to!

Summary of borrowing rules

- References are always valid and non-null.
- Every value has one owner.
- Values are freed when their owner goes out of scope.
- Assignment moves values (unless the value is Copy).
- Values that allocate memory on the heap are usually not Copy.
- You can have one mutable reference, or multiple immutable references. But not both.
- A reference can never outlive its value.

- Problems ...

Problem 1

```
fn main() {  
    let mut v = vec![5, 4, 3, 2];  
    append_one(&v);  
}  
  
fn append_one(v: &mut Vec<i32>) {  
    v.push(1);  
}
```

- (The `vec!` macro just initializes a `Vector`, which is a dynamically sized array-based list. A `Vec` is not `Copy`.)
- Incorrect types: we're passing an *immutable* reference to a function that expects a *mutable* reference.

Problem 2

```
fn main() {  
    let v = vec![5, 4, 3, 2];  
    append_one(v);  
    assert_eq!(v[4], 1);  
}  
  
fn append_one(mut v: Vec<i32>) {  
    v.push(1);  
}
```

- Use of moved value: `v` is moved when we call `append_one`, so we're not allowed to use it in the `assert_eq` statement.

Problem 3

```
fn main() {  
    let mut v = vec![5, 4, 3, 2];  
    let mut v = append_one(v);  
    assert_eq!(v[4], 1);  
}  
  
fn append_one(mut v: Vec<i32>) -> Vec<i32> {  
    v.push(1);  
    v  
}
```

- No problems here! `v` is moved into `append_one`, but `append_one` returns `v`. So `v` is moved back into the caller.

Problem 4

```
struct Rect {  
    width: u32,  
    height: u32,  
}  
  
impl Rect {  
    fn transpose(&mut self) {  
        let tmp = self.width;  
        self.width = self.height;  
        self.height = tmp;  
    }  
}  
  
fn main() {  
    let r = Rect { width: 2, height: 5 };  
    let ptr = &r;  
    r.transpose();  
    assert_eq(ptr.width, 5);  
}
```

- `r` is mutably borrowed (in `transpose`) while immutably borrowed (to `ptr`)! Not allowed.

Problem 5

```
fn main() {  
    let v = vec![1, 2, 3, 4];  
    let one = &vec[0];  
    vec.push(5);  
    println!("1 = {}", *one);  
}
```

- Cannot borrow `vec` mutably (to push) while it is already immutably (to `one`) borrowed.
- Think about this: the call to push might result in the `vec` being realloc'd. This might invalidate the `one` pointer. borrowed while immutably borrowed! Not allowed.

Convenience types

- Dynamically sized arrays.
- Create a Vec:

```
// The type annotation is needed if the compiler can't determine the element type.  
let x: Vec<i32> = Vec::new();
```

- Or use the vec! macro:

```
let x = vec![1, 2, 3];  
let y = vec![162; 3]; // Equivalent to vec![162, 162, 162].
```

- Operations on a Vec:

```
let mut x = vec![1, 2, 3];  
assert_eq(x.len(), 3);  
x.push(4);
```


- Pointers are never null! What if you actually *want* something to be null?
- Use an `Option<T>`! Here's the definition of `Option`, from the standard library:

```
pub enum Option<T> {  
    None,  
    Some(T),  
}
```

- Two possible cases: the option is either `None`, or it is `Some`.
- If an `Option` is `Some`, the value in the `Some` variant will always be a valid value of type `T`.
- The `T` is a generic type parameter. It behaves like generics in other languages, such as Java.
- This prevents us from having to manually define separate `Option` types for `i32`, `String`, etc.

■ Here's how you can use an option:

```
// This type annotation is not necessary.  
let x: Option<i32> = Some(4);  
assert!(x.is_some());  
let y = x.unwrap(); // get the value out of Option  
assert_eq!(y, 4);  
  
// This type annotation IS necessary!  
let z: Option<i32> = None;  
assert!(z.is_none());  
  
let w = Some(String::from("hello"));  
match w {  
    Some(s) => println!("{}", world!", s),  
    None => panic!("didn't expect to get here"),  
}
```

- Most languages handle errors via one of two ways:
 - Try/catch exceptions (Python, Java, JavaScript, etc.)
 - Returning a separate error value (Go)
 - Many times, errors are not handled properly or are tedious to handle.
- Rust tries to make error handling easier. It's not always smooth sailing though.

- The idiomatic way to handle errors in Rust is via `Result<T, E>`:

```
pub enum Result<T, E> {  
    Ok(T),  
    Err(E),  
}
```

- Functions that may not complete successfully should return a `Result`.
- If the result is `Ok`, the caller can access the returned value (of generic type `T`).
- If the result is `Err`, additional information (of generic type `E`) about the error is returned.
- If a program cannot recover from an error, you can `panic!` instead of returning a `Result`. The panic will immediately terminate the program.

■ Example:

```
use std::fs::File;

fn main() {
    let f = File::open("hello.txt");
    let f = match f {
        Ok(file) => file,
        Err(error) => panic!("Problem opening the file: {:?}", error),
    };
}
```

- Opening a file is **fallible**, so `File::open` returns a `Result`. We check if the open was successful; if not, we panic and exit.

- Matching on Results all the time can be tedious. If we know we are going to panic on an Err, we can use `unwrap()` instead:

```
use std::fs::File;

fn main() {
    let f = File::open("hello.txt").unwrap();
}
```

- `unwrap()` panics on error; otherwise, it returns the data contained in the Ok variant.

- Another shortcut is the `?` operator:

```
use std::fs::File; use std::io; use std::io::Read;

fn read_file() -> Result<String, io::Error> {
    let mut f = File::open("hello.txt"?);
    let mut s = String::new();
    f.read_to_string(&mut s)?;
    Ok(s)
}
```

- If a `Result` is an `Err`, `?` causes the function to return that `Err`. Otherwise, `?` unwraps the `Ok` variant.
- It's actually a bit more complicated than that. The `?` also attempts to do some type conversion: If your function returns `Result<T, E1>`, but you apply `?` to a `Result<U, E2>`, Rust will try to convert error `E2` into an error of type `E1`.

- The Rust book

- <https://doc.rust-lang.org/stable/book/>

- Rust by Example

- <https://doc.rust-lang.org/stable/rust-by-example/>

- Smart pointers

- <https://doc.rust-lang.org/book/ch15-00-smart-pointers.html>

- Generics, Traits, and Lifetimes

- <https://doc.rust-lang.org/book/ch10-00-generics.html>